

Hyperparameter Tuning for Milvus via HOB0

Advisor: Zhao Lyu
Name: Xiang Pan
September 29, 2021

Contents

1	Introduction	2
2	Methods	2
2.1	BOHB	2
2.2	Loss Function	2
3	Implemtation	3
3.1	Hardware Information	3
3.2	Core Code Annotations	3
3.2.1	ENV	3
3.2.2	Search Config Class	5
3.2.3	ENV Input	5
3.3	User Input Parser	5
3.4	HOBO	6
4	Results	6
5	Attemption	8
5.1	VAE	8
5.1.1	Measure of Rank Keeping	8
5.1.2	Methods	8
5.2	Graph Based Message Passing	9
5.3	Minimum Distortion Embedding	9

1 Introduction

- Project Name: Auto tuner for vector indexing parameters(210310187)
- Scheme Description: Auto tuner for milvus and vectors' preprocessing for milvus' friendly vectors space

To best serve users' demand of performance on Milvus, we use the Bayesian Optimization and Hyperband(BOHB)[2] as our parameter search method to optimize Milvus parameters on three stages: the index-type, the index build, and the index search.

2 Methods

To get an end-to-end solution for index choices, we apply the BOHB in different levels. Notably, the index-type plays a crucial role in deciding the final performance on Milvus given datasets. However, Bayesian Optimization(BO) may not fully explore some specific types due to initially poor performance. This means the randomness of Bayesian Optimization may lead the searching process to fall into a local optima. Therefore, we set two index type optimization modes: the direct BO and a loop over the index types.

2.1 BOHB

Please see Fig. 1.

2.2 Loss Function

We use Laplace's method to convert a constrained BO to an unconstrained version. Our loss function is defined as below:

$$Loss = sign(recall, threshold) - query_per_sec \quad (1)$$

$$Sign(recall, threshold) = \begin{cases} recall - threshold & recall > threshold \\ \lambda \cdot (threshold - x) & recall \leq threshold, \end{cases} \quad (2)$$

Here we set $\lambda = 100000$ for Lagrange method, and $threshold = 95$.

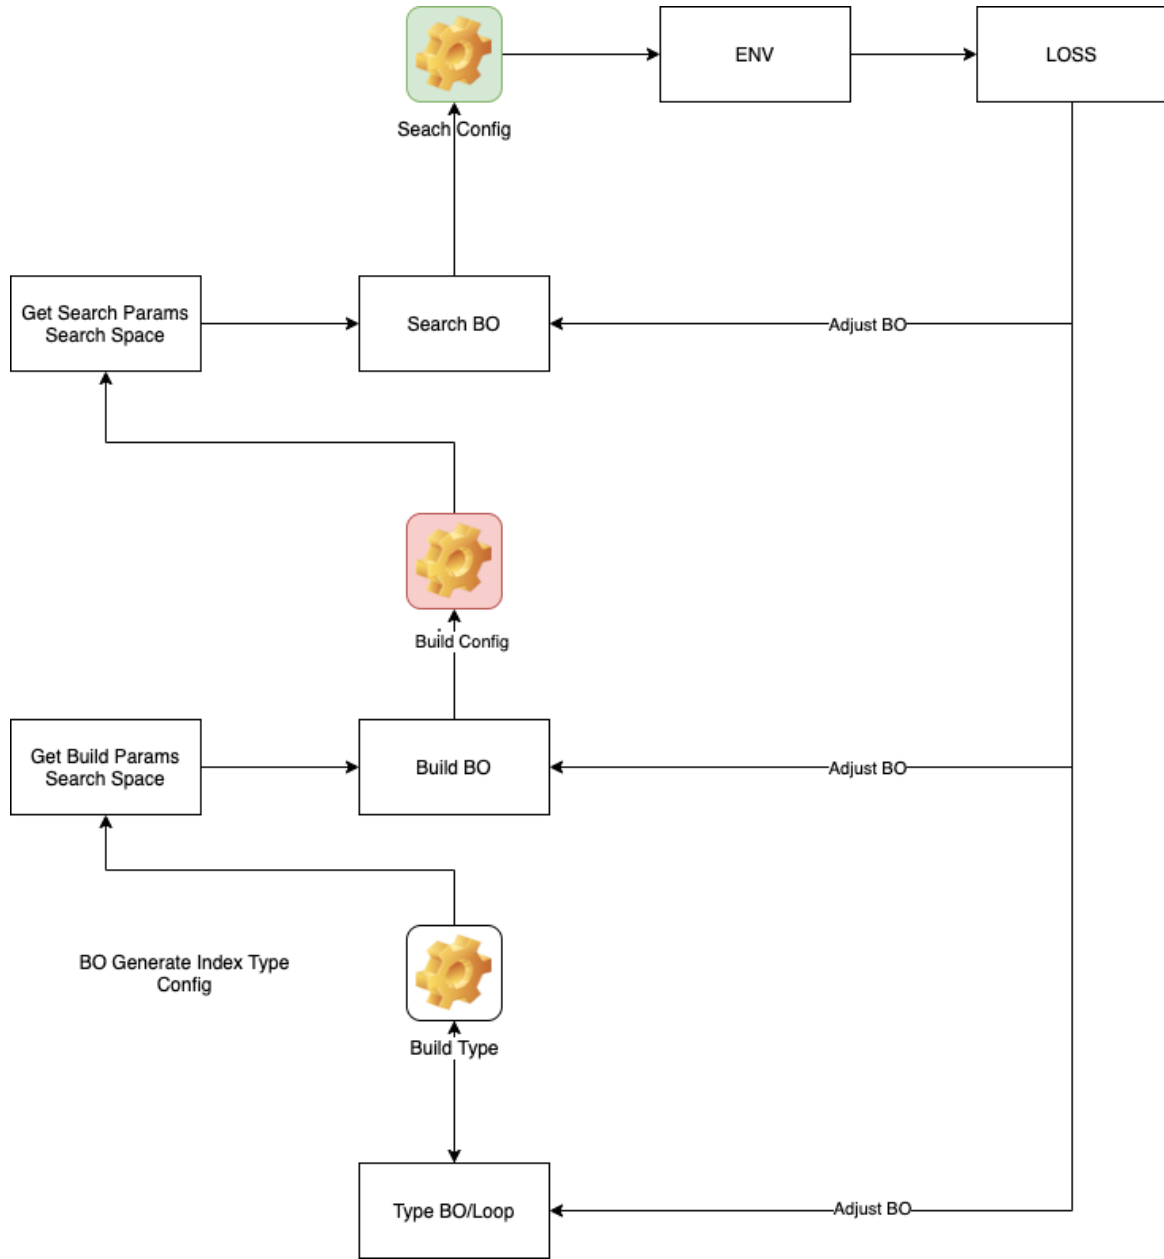


Figure 1: Flow Diagram

3 Implemtation

3.1 Hardware Information

Our method is tested on server with: CPU: Intel Core i7-8700 CPU @ 4.6GHz and RAM: 32083MiB.

3.2 Core Code Annotations

3.2.1 ENV

ENV is a helper class which configures basic Milvus related variables.

Listing 1: ENV class config

```

1 class ENV():
2 def __init__(self, args = None):
3     print("ENV")
4     # docker related information
5     host = '127.0.0.1'
6     port = '19530'
7
8     # get milvus client and collection_name
9     self.client = Milvus(host, port)
10    self.collection_name = args.collection_name
11
12    # get query_vectors and set top_k
13    self.query_groundtruth = self.get_groundtruth()
14    self.query_vectors = self.get_query()
15    self.top_k = 100
16
17    # get status by current db
18    self.index_type = None
19    self.index_params = None
20    self.refresh_status()
21
22    # set datadim, which is needed by some search constraint
23    global gDataDim
24    gDataDim = 128
25
26    # based on the input type, get the default build config
27    if args.op == "build_params":
28        self.target_index_type = get_index_type(args.index_type)
29        self.target_index_params = \
30            get_default_build_config(self.target_index_type)
31
32        is_build = False
33        if self.index_type != self.target_index_type:
34            is_build = True
35        elif self.index_params != self.target_index_params:
36            is_build = True
37
38        if is_build:
39            self.env_build_input(self.target_index_type, \
40                                self.target_index_params)
41            self.refresh_status()
42
43    # set search space
44    self.default_build_config = get_default_build_config(self.index_type)
45    self.search_configspace = get_search_configspace(self.index_type, \
46                                                       self.index_params)
47    self.build_configspace = get_build_configspace(self.index_type)

```

Listing 2: Refresh Status

```

1 # base on user's input, get the target search space
2 def refresh_status(self):
3     """
4     refresh status
5     reset index_type and index_params
6     reset all config space
7     """
8     status, stats = self.client.get_index_info(self.collection_name)
9     self.index_type = stats._index_type

```

```

10     self.index_params = stats._params
11
12     # set config space
13     self.default_build_config = get_default_build_config(self.index_type)
14     self.build_configspace = get_build_configspace(self.index_type)
15     self.search_configspace = get_search_configspace(self.index_type, \
16         self.index_params)

```

3.2.2 Search Config Class

Listing 3: Search Config Class

```

1 class HNSW_build_search_shared_config(object):
2     def __init__(self):
3         self.M = cs.IntegerUniformHyperparameter('M', 4, 64)
4         self.efConstruction = \
5             cs.IntegerUniformHyperparameter('efConstruction', 8, 512)
6         top_k = 100
7         self.ef = cs.IntegerUniformHyperparameter('ef', top_k, 512)
8         self.configspace = \
9             cs.ConfigurationSpace([self.M, self.efConstruction, self.ef], seed=123)

```

We use constant class to configure the default search space.

3.2.3 ENV Input

Listing 4: Input

```

1 # given full env put
2 def config_input(self, config):
3     # check current index type
4     is_build = False
5
6     self.refresh_status()
7     if self.index_type != config['index_type']:
8         is_build = True
9     elif self.index_params != config['index_params']:
10         is_build = True
11
12     if is_build:
13         self.env_build_input(config['index_type'], config['index_params'])
14         self.refresh_status()
15
16
17     recall, query_per_sec = self.env_search_input(config['search_params'])
18     self.search_params = config['search_params']
19
20     return recall, query_per_sec

```

For input, we check the current status and mark the change that needs to be done. Once we have changed the index of current database, we refresh the ENV status.

3.3 User Input Parser

Listing 5: Input Parser

```

1  if args.op == "build_type":
2      build_type_search_spcas = \
3          [IndexType.IVF_FLAT, IndexType.IVF_PQ, IndexType.IVF_SQ8, IndexType.HNSW]
4      if args.build_type_op_method == "BO": # BO
5          index_type = \
6              cs.CategoricalHyperparameter('index_type', build_type_search_spcas)
7          index_type_configspace = cs.ConfigurationSpace([index_type], seed=123)
8          type_opt = \
9              BOHB(index_type_configspace, \
10                  build_type_evaluate, max_budget=10, min_budget=1)
11         type_logs = type_opt.optimize()
12     else: # Loop
13         for index_type in build_type_search_spcas:
14             env.target_index_type = index_type
15             env.refresh_status()
16             opt = \
17                 BOHB(get_build_configspace(env.target_index_type), \
18                     build_evaluate, max_budget=10, min_budget=1)
19             logs = opt.optimize()

```

The input parser is trivial.

3.4 HOB0

Listing 6: HOB0

```

1  # Based on the HOB0 class and our search space configuration, \
2  # we can build a BOHB object. After that, we can call the optimize() \
3  # method to start the optimization process.
4  def build_type_evaluate(params, n_iterations):
5      env.target_index_type = params['index_type']
6      env.refresh_status()
7      if args.build_search_share_space:
8          opt = BOHB(get_build_search_shared_configspace(env.target_index_type),
9                     build_search_share_space_evaluate,
10                     max_budget=n_iterations,
11                     min_budget=1,
12                     eta = 10)
13         logs = opt.optimize()
14     else:
15         opt = BOHB(get_build_configspace(env.target_index_type),
16                    build_evaluate,
17                    max_budget=n_iterations,
18                    min_budget=1,
19                    eta = 10)
20         logs = opt.optimize()
21     return logs.best['loss']

```

4 Results

HOBO Results

Set 25, 2021

Results

Index Type Optimization

Method	index_type	M	efConstruction	ef	recall	query_per_sec	loss
BOHB(Index Type Loop)	'HNSW'	17	445	114	99.68	16782.6	-16777.9
BOHB(Index Type BO)	'HNSW'	22	274	106	99.75	18522	-18517.2

	index_type	nlist	M	nprobe	recall	query_per_sec	loss
Grid Search	'HNSW'	4	158	200	97.11	18331.74	-18329.63

Index Parameters Optimization

IVF_FLAT

	index_type	nlist	nprobe	recall	query_per_sec	loss
BOHB	'IVF_FLAT'	2883	54	99.68	14911	-14906.3
Grid Search	'IVF_FLAT'	14601	101	100.0	14402.03	-14397.03

IVF_SQ8

	index_type	nlist	nprobe	recall	query_per_sec	loss
BOHB	'IVF_SQ8'	8405	46	98.86	13827.5	-13823.7
Grid Search	'IVF_SQ8'	5401	101	99.49	13080.62	-13076.13

IVF_PQ

	index_type	m	nlist	nprobe	recall	query_per_sec	loss
BOHB	'IVF_PQ'	128	3800	205	98.1	1289.00	-1285.90
Grid Search	'IVF_PQ'	64	1	1	95.08	1733.67	-7629.25

HNSW

Method	index_type	M	efConstruction	ef	recall	query_per_sec	loss
BOHB	'HNSW'	18	92	157	99.85	17868.6	-17863.8
Grid Search	'HNSW'	4	158	200	97.11	18331.74	-18329.63

5 Attention

5.1 VAE

We tried using VAE to preprocess the given datasets to compress the embedding dimension and get a more Milvus-friendly representation space.

5.1.1 Measure of Rank Keeping

We define the overlap and exact-match to measure the rank-keeping performance.

$$overlap = \frac{1}{|D|} \sum_{eb \in D} \frac{knn \text{ of } f(eb) \cap knn \text{ of } eb}{knn \text{ of } eb}, \quad (3)$$

$$exact \text{ match} = \frac{1}{|D|} \sum_{eb \in D} \frac{\sum_{neb \in neighbor \text{ of } eb} rank \ f(neb) == rank \text{ of } neb}{k}, \quad (4)$$

where D represents the given dataset, eb is an vector in the dataset, and f denotes the encoder function.

5.1.2 Methods

Minimum Reconstruct Error

$$Loss_{MRE}(eb) = decoder(encoder(eb)) - eb \quad (5)$$

We use $compressed_eb = encoder(eb)$ for milvus search.

Distance Preserving

$$Loss_{DP}(e1, e2) = \|e1 - e2\|_2 - \|encoder(e1) - encoder(e2)\|_2 \quad (6)$$

Contrastive Learning The key idea of contrastive learning is to use contrastive loss to have the embedding remain the relative distance, which "may" be useful to keep the rank.

$$Loss_{CL}(a, p, n) = \|a - p\|_2 - \|a - n\|_2, \quad (7)$$

where a is the anchor, p is the positive/similar and n is the negative/dissimilar. Positive and negative attributes can be decided by their relative distance to a . Similarly, we have,

$$Loss_{CL}(a, p, n, f) = \|f(a) - f(p)\|_2 - \|f(a) - f(n)\|_2, \quad (8)$$

f is the encoder function.

- If p is closer to a than n , then p is positive and n is negative.
- If p is top- k nearest neighbors of a but n is not, then p is positive and n is negative.

Our empirical result shows that the loss of contrastive learning and VAE can not preserve the relative distance rank between embeddings.

5.2 Graph Based Message Passing

We use the message passing to make the more similar embedding closer. First, we build the knn graph g for given dataset D . Over graph g , we define the following message passing function,

$$u' = u + \text{reduction}(\text{aggregation}(v, e)) \quad (9)$$

For each node u in g , v represents its neighbors, e denotes a weighted edge between u and v , and the aggregation is the aggregation function.

We use $\text{aggregation} = e_mul_v$ and $\text{reduction} \in \{\text{mean}, \text{max}, \text{min}, \text{sum}\}$. The result shows that max get the best performance, which may suggest that simply average the neighbors' embedding is not enough. However, the best overlap performance is around 0.65, which we do not think is good enough for the ranking preserving.

5.3 Minimum Distortion Embedding

We applied the method proposed in [1], which is to minimize the distortion of the embedding rank while compressing the embedding dimension.

References

- [1] Akshay Agrawal, Alnur Ali, and Stephen Boyd. Minimum-Distortion Embedding. page 179.
- [2] Stefan Falkner, Aaron Klein, and Frank Hutter. BOHB: Robust and Efficient Hyperparameter Optimization at Scale. page 10.